

© LIVE DEMO · NO SIGNUP · NO CARD
NO. 009 — SEE AN AGENT FINISH THE JOB

Watch an AI agent automate a workflow — live.

Press start. We will simulate one of our healthcare-intake agents processing a real (anonymised) incoming fax. 12 seconds. No signup required. Pick a different workflow from the left if you'd rather see legal, finance, or hiring.

SCENARIO

STEP 1 OF 2

01 · DEFAULT

1

HEALTHCARE INTAKE

PROCESS AN INCOMING PATIENT FAX: EXTRACT FIELDS, CROSS-CHECK THE EMR, FLAG EXCEPTIONS.

02

1

LEGAL DISCOVERY

TRIAGE, TAG, AND PRIVILEGE-FLAG A FRESHLY UPLOADED LITIGATION EXHIBIT.

03

1

FINANCE RECONCILIATION

MATCH A BANK STATEMENT LINE TO THE CORRECT GL ENTRY AND POST THE JOURNAL.

HR · CV SHORTLIST

READ A CV, SCORE IT AGAINST THE JOB SPEC, AND DRAFT A SHORTLIST NOTE FOR THE HIRING MANAGER.

Note. Inputs are anonymised samples from real production traffic – the agent logic that runs against them is the exact same code we ship to clients.

AGENT LOG · LIVE STREAM

● IDLE

○ ○ ● autopilotlabs@ops ~ /agents/intake.01/run --scenario=healthcare

SECURE · TLS 1.3



Press **start demo** to watch the agent parse, validate, and ship an outcome.
Takes about 12 seconds.

NO SIGNUP · ANONYMISED INPUT · 1 OUTCOME PER RUN

▶ Start demo

↺ Reset

METRICS

LIVE

OUTCOMES SHIPPED

+1

0

THIS SESSION



AVG LATENCY

—

0ms

FAX RECEIVED → OUTCOME



COST · PER OUTCOME

—

A\$0.00

VS A\$18 HUMAN BASELINE



CONFIDENCE

—

0.00

AGENT SELF-ASSESSED



◆ Local session

Run 00

— § 02 — WHAT JUST HAPPENED

Three phases. **Twelve seconds.** One finished job.

What you watched is the same three-phase loop every Autopilot Labs agent runs. No chatbot, no prompt engineering by you — the AI reads the input, checks itself against your system of record, then ships the outcome to wherever it needs to go. In plain English, here is what it does, and why it matters.

— 01 · PARSE



Read the input, field by field.

The agent opens the incoming document — a fax, an email attachment, a bank CSV, a CV — and pulls out structured fields. Names, dates, amounts, references. Built-in OCR, vision, and table extraction are all handled by the AI. Nothing for you to configure.

IN PLAIN ENGLISH

Think of a very fast, very careful intern who can read anything, in any layout, in about a second — and never mistypes.

— 02 · VALIDATE



Cross-check against your system of record.

The AI does not just guess. It checks every extracted field against your source of truth — the EMR, the ledger, the ATS, the warehouse. If anything disagrees, it flags the exception, scores its own confidence, and routes low-confidence items to a human. The rest, it finishes.

IN PLAIN ENGLISH

Before the agent ships anything, it asks: "Do I agree with this? Does my database agree with this? Am I sure?" If no, a human sees it. If yes, it ships.

— 03 · SHIP



Post the outcome where it belongs.

The finished job — a normalised record, a posted journal entry, a tagged document, a draft email — lands in your existing system. No new UI for your team to learn. The agent logs every action, timestamps the outcome, and bills only if the SLA passed. Miss the SLA, you do not pay.

IN PLAIN ENGLISH

The work shows up in the tool your team already uses — Epic, NetSuite, Relativity, your ATS — as if a very diligent teammate did it while you were asleep.

§ 03 — STOP WATCHING. START SHIPPING.

Want one of these automating your workflow?

We will spend ten working days inside your team, map the path end-to-end, and come back with a shortlist of workflows we can automate with AI — priced per finished outcome. If nothing fits, you walk away with a free diagnostic report.

Book a free audit



See outcome pricing



Browse our agents



About the [demo](#), in plain English.

§ 04 — FAQ

Is this the actual AI running, or a recording?

The log stream is scripted for the demo so you see the full story in 12 seconds. The phases, fields, and timings mirror a real production run. Want to watch the real one? Book an audit and we will give you a read-only seat on the ops console.

Do you see my data?

No. Every byte on this page runs locally in your browser. Nothing is sent to us, logged, or stored. Pick a scenario, press start, close the tab. We never know you were here unless you book.

What if I want a different workflow simulated?

The four scenarios cover our most common agent shapes. If yours is different — a claim triage, a lease abstraction, a return dispute — we will build a bespoke demo on the audit call, with your own sample data.

How accurate are the latency and cost figures?

They match the 90-day rolling medians from our production ops console. Healthcare-intake actually runs a hair faster; legal discovery a hair slower on very large exhibits. Everything within a couple of hundred milliseconds.

Can the AI be wrong? What happens then?

Yes — and it is built to know when. Every agent scores its own confidence. Below a per-workflow threshold, the outcome is routed to a human reviewer before it ships. You never pay for an outcome that did not pass the SLA.

How soon could something like this run on my workflow?

Median audit-to-production is 14 days after contract. The audit itself is 10 working days, so from first call to finished agent you are looking at about a month — no new hires, no new SaaS, no new training.

Agents that finish.

Autopilot Labs is an AI workflow studio with hubs in Melbourne, Dubai and Riyadh. We build vertical agents across fourteen practice areas — thirteen industries plus an always-on voice desk in 14 languages (Arabic included) — and bill per completed outcome. Not per seat, not per token.

COMPANY

→ [Solutions](#)

→ [Process](#)

→ [Pricing](#)

→ [Case studies](#)

→ [About](#)

→ [Contact](#)

VERTICALS

→ [Healthcare](#)

→ [Legal](#)

→ [Finance](#)

→ [HR & hiring](#)

OFFICE

→ [L14 · 459 Collins St](#)

→ [Melbourne VIC 3000](#)

→ hello@autopilotlabs.io